

Yashraj Parmar

Data Glacier

2 June 2025

Week 9: Deliverables

1) Information:

- a) Group Name: NLP Sentinels
- b) Name: Yashraj Parmar
- c) Email: yparmar2024@gmail.com
- d) Country: United States of America
- e) College: Stevens Institute of Technology
- f) Specialization: Natural Language Processing
- g) Internship Batch: LISUM44

2) Problem Description:

- a) Social media platforms rely on hate speech detection systems to monitor and filter content that may contribute to online harm, including cyberbullying. However, these systems are not always fully effective, and harmful content can still go undetected. The goal of this project is to develop a machine learning model using natural language processing (NLP) techniques that can accurately assess and classify the level of hate speech in a given piece of text.

3) Data Cleansing:

- a) After checking for null values in the dataset, I realized that there were no null or missing values in any of the columns. Therefore, nothing had to be done for the null/missing values.

- b) Since the dataset had only two possible values: 0, or non-hate speech, and 1, hate speech, outlier detection was not applicable in this case.
- c) The class imbalance of significantly more non-hate speech examples than hate speech examples is addressed through the use of the `class_weight = "balanced"` parameter which allows the model training phase to ensure fair learning from both classes.
- d) Most of the text has inconsistency in casing, the whole tweet can be reduced to lowercase to avoid treating same words with different casing as separate tokens. This can be removed through Python's `.lower()` function which lowers the case of all contents of a string.
- e) The text has user mentions such as "`@user`" which had no relevant information to the sentiment of the text, therefore this can be removed as well. This can be removed through Python's `.replace()` function which replaces all instances of "`@user`" with "".
- f) The text contains hashtags, and while these contain sentimental value, the "`#`" symbol is irrelevant, while the word is useful therefore the word will stay. This can be removed using Python's `.replace()` function as well which will replace all instances of "`#`" with "".
- g) The text contains numerous special characters which will only cause noise in the model's interpretation, therefore this can be removed as well using Python's Regex library. This can be removed using Python's `re.sub()` function which will replace everything except lowercase letters, numbers, and spaces.

- h) Some of the text also contains trailing spaces occasionally which is standardized spacing and removes leading/trailing whitespace. This can be removed using Python's `re.sub()` function as well which will replace all white spaces with nothing.

4) Transformation Done on Data:

- a) The first featurization technique I will be using is TF-IDF Vectorization from scikit-learn. This method transforms text into numerical vectors based on the frequency of words and their importance across documents (Term Frequency–Inverse Document Frequency). Common words across all texts receive lower scores, while unique or meaningful words get higher importance. A summary of the steps to take towards this are: tokenize text data and lowercase it, remove stop words, generate a sparse matrix of TF-IDF scores where each row represents a tweet and each column a word feature, and finally the train the representation using a baseline logistic regression model. This featurization's purpose is to provide a traditional NLP baseline and compare its performance to transformer-based approaches.
- b) The second featurization technique I will be using is BERT Embeddings, the `bert-base-uncased` model from Hugging Face Transformers. This method uses a pre-trained BERT model to convert each tweet into a dense 768-dimensional vector representation that captures semantic meaning and context. A summary of the steps to take towards this are: to clean tweets that were tokenized using BERT's tokenizer, pad the sequences and truncate it to a fixed length, extract the [CLS] token embedding from the final hidden state to represent the entire tweet,

use these embeddings to train a deep learning classifier such as a feedforward neural net. This featurization's purpose is to leverage deep contextual representations of text that outperform traditional vectorizers on semantic tasks like hate speech detection.